

DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING
INDIAN INSTITUTE OF TECHNOLOGY MADRAS

Computer Organization and Architecture (CS2600)

Total Time : 50 minutes
Max Marks : 25

Quiz 2
20-03-2024

-
1. Answer all questions
 2. Unless otherwise specified, all questions are related to RISC V architecture.
 3. State clearly any assumptions that you make.
 4. 1-mark for neatness.
-

1. Consider a 32-bit CPU with a direct-mapped, byte-addressable write-through cache memory of 32 KByte, having a cache line size of 64 bytes. (9 marks)

- (a) How many cache lines are present in this cache memory? (1 mark)

Answer: Cache lines = $32KB/64 = 2^{15-6} = 2^9 = 512$

Grading: 0 or 1

- (b) What is the length of the tag bits in the 32-bit address? (1 mark)

#bits in tag per cache line = $32 - (9 + 6) = 17$

Grading: 0 or 1

- (c) What are the total number of bits present in the cache? (2 marks)

Answer: #bits of data in cache line = $64 \times 8 = 2^{6+3} = 2^9$

#valid bit = 1

#Total number of bits = $512 \times (512 + 17 + 1) = 2,71,360$ bits

Grading: 1 for the formula; 1 mark for the answer

- (d) A program defines an array as follows `int T0[1024]` and then immediately reads each element of the array as follows:

```
for(int i=0; i<1024; ++i) x ^= T0[i];
```

If a cache hit takes 1ns and a cache miss takes 10ns to complete a load/store operation, what is the minimum time the program takes to access the entire array? (3 marks)

Answer: #Number of elements per cache line = $64/\text{sizeof(int)} = 64/4 = 16$

Minimum execution time occurs when the array is aligned to 64 bytes.

#Number of cache misses to read the entire array into cache is $\text{sizeof(T0)}/16 = 4096/16 = 64$

#Thus, we have 64 cache misses and the remaining $1024 - 64 = 960$ cache hits. Total memory access time is $64 \times 10 + 960 \times 1 = 1600\text{ns}$

Grading: 1 mark for number of cache misses; 1 mark for cache hits; 1 mark for the answer.

- (e) The program then executes the function:

```
memset(T0, 0, sizeof(T0));
```

which sets all elements of T0 to 0. How much time does the `memset` function take to execute memory operations? Assume no write buffer is present. (1 mark)

Answer: Every store would result in a write back to RAM. Thus the `memset` would take $1024 \times 10 = 10240\text{ns}$.

Grading: 0 or 1

- (f) How would a write buffer improve the performance of the processor? (1 mark)

Answer: With a write buffer, the CPU does not stall while the memory block updates the memory. This boosts performance. <(optional: For example, ideally, the entire `memset` function executes in 1024ns.)>

Grading: part marks possible

2. Consider a 32-bit CPU with a 2-level paging unit with each page table entry is of 4Byte, and each page is of 4KByte. (8 marks)

- (a) How many bits of address are needed to index into a second-level page table? (1 mark) Number of bits to index into a second level page is $4KB/4 = 1024 \rightarrow 10\text{-bits}$

- (b) What is the maximum amount of memory (in bytes) needed to store the page tables for a single program? (1 mark)

First level page table has 2^{10} entries. Thus, there are 2^{10} second level pages. The total memory required to store all page tables is $4096 \times (2^{10} + 1)$

(c) A program defines an array as follows `int T0[16384]` and then reads each element of the array as follows:

```
for(int i=0; i<15360; ++i) x ^= T0[i] ^ T0[i+1024];
```

- i. What is the maximum number of memory accesses needed for the array accesses during the execution of the `for` loop?
 Each iteration makes two accesses to the array. If there is in no TLB, each access would require 3 memory operations. Two for converting from virtual address to physical address, the third for the actual memory operation. Thus, the number of memory accesses is $3 \times 2 \times 15360 = 92,160$
 Grading: 2 marks. 1/2 mark for TLB assumption. 1/2 mark for 3 memory operations, 1 mark for calculation.
- ii. What is the maximum number of these memory accesses that can result in page faults? Its a 2-level page table. 2 page faults to load the page directory and a second level page table. Every element of the array is loaded. There are 16384 elements. If the array is aligned to 4KB, it spans across 16 pages. Thus, a maximum of $16+2=18$ page faults are obtained. If the array is not-aligned to 4KB, it spans across 17 pages. Thus, a maximum of $17+2=19$ page faults are obtained.
 Grading: 3 marks for 1 mark for page dir and page table faults; 2 mark for non-alignment; 1 mark if only alignment is mentioned;

(5(2+3) marks)

(d) Suppose the CPU designer wants to introduce a TLB in the processor. What is the minimum size of the TLB that best optimizes the `for` loop described above? Ignore all other memory accesses. (1 marks)
 At any given iteration, the program accesses at-most 2 pages. Thus, a TLB of size 2 would provide optimal result.

3. Match the elements in Table 1, based on what the LHS is **most related to**. Each entry on the LHS uniquely matches an entry in the RHS. (5 marks)

Table 1: Match the Following	
(A) mpp	(a) is most related where a program is interrupted
(B) mepc	(b) is most related to privilege modes
(C) mtvec	(c) is most related to the source of the interrupt
(D) mcause	(d) most related to interrupt service routines
(E) mret	(e) most related to resuming execution

- (A) → (b)
- (B) → (a)
- (C) → (d)
- (D) → (c)
- (E) → (e)

4. An OS has a system call defined as follows:

```
int my_read(int fd, char *buf, int nbytes);
```

Write a RISC V assembly code snippet to explain how this system call is realized in assembly in the user program and explain. (3 marks)
 0.5 marks for fd, buf, nbytes in a0, a1, a2 resp. 0.5 marks for a7 holding system call number; 0.5 marks for ecall; 0.5 marks for return; 1 marks for clarity in explaining.